

HAN AEA - Embedded Vision & Machine Learning

EVD1 - Week 2

Image Fundamentals Graphics Algorithms

By Hugo Arends

Image Fundamentals

- Functions for creating and deleting images
- Functions for converting images
- Functions for reading and writing pixels
- Basic image processing operators
- Scaling fast



Scaling

- Used to enhance contrast
- Used to scale larger pixel datatypes to smaller pixel datatypes
e.g. float to basic
- Min-max scaling is defined as

$$p_{dst}(x, y) = \frac{dst_{max} - dst_{min}}{src_{max} - src_{min}} \cdot (p_{src}(x, y) - src_{min}) + dst_{min}$$

Scaling - example

$$p_{dst}(x, y) = \frac{255 - 0}{src_{max} - src_{min}} \cdot (p_{src}(x, y) - src_{min}) + 0$$



Scaling operation

$$p_{dst}(x, y) = \frac{dst_{max} - dst_{min}}{src_{max} - src_{min}} \cdot (p_{src}(x, y) - src_{min}) + dst_{min}$$

Requires a lot of operations for every destination pixel:

- Read all source pixels to determine src_{max} and src_{min}
- Calculate stretch factor $\frac{dst_{max}-dst_{min}}{src_{max}-src_{min}}$
- Calculate new pixel value $p_{dst}(x, y)$ (with a point number) and store the result

Performance optimization

- Optimize the implemented code for execution speed
- Several techniques discussed in random order

Performance optimization

Use the compiler optimization levels

- None (-O0)
- Optimize (-O1)
- Optimize more (-O2)
- Optimize most (-O3)
- Optimize for size (-Os)
- Optimize for Debug (-Og)


Performance optimization

Use the compiler optimization levels

For the given project there are two build configurations, making it easy to switch between optimization levels

- Optimize for debug (-Og)
- Optimize most (-O3)

TIP. Change the build configuration in the Projects pane



Performance optimization

- Optimize for debug (-Og)

~850 µs

- Optimize most (-O3)

~230 µs

```
// Copy image
uint8_pixel_t *s = (uint8_pixel_t *)src->data;
uint8_pixel_t *d = (uint8_pixel_t *)dst->data;

dst->rows = src->rows;
dst->cols = src->cols;
dst->type = src->type;

for(int32_t r = src->rows-1; r >= 0; r--)
{
    for(int32_t c = src->cols-1; c >= 0; c--)
    {
        *d++ = *s++;
    }
}
```

Performance optimization

Use the built in FPU!

- Is already enabled in the given project, because it is required by the video driver
- Will you be using ‘doubles’ or ‘floats’?


```
src->data[0] += 0.5; // Is 0.5 a float or a double?  
src->data[0] += 0.5f; // Is 0.5 a float or a double?
```

- Alternative (if your hardware doesn’t support a hardware FPU): implement Fixed Point calculations

Performance optimization

Use register variables

- *Register variables require less memory access*
- *Especially useful for loop-counters, because these are accessed often*
- *Although taken care of by the compiler, you can explicitly use the register keyword*
- [https://en.wikipedia.org/wiki/Register_\(keyword\)](https://en.wikipedia.org/wiki/Register_(keyword))



Source: MCXNx4x datasheet

Performance optimization

Use pointers to the source and destination data instead of getter and setter functions

- Optimize most (-O3)

$\sim 230 \mu s$

```
// Copy image
uint8_pixel_t *s = (uint8_pixel_t *)src->data;
uint8_pixel_t *d = (uint8_pixel_t *)dst->data;

dst->rows = src->rows;
dst->cols = src->cols;
dst->type = src->type;

for(int32_t r = src->rows-1; r >= 0; r--)
{
    for(int32_t c = src->cols-1; c >= 0; c--)
    {
        *d++ = *s++;
    }
}
```

Performance optimization

Use pointers to the source and destination data instead of getter and setter functions

- Optimize most (-O3)

~4300 µs

```
// Copy image
uint8_pixel_t *s = (uint8_pixel_t *)src->data;
uint8_pixel_t *d = (uint8_pixel_t *)dst->data;

dst->rows = src->rows;
dst->cols = src->cols;
dst->type = src->type;

for(int32_t r = src->rows-1; r >= 0; r--)
{
    for(int32_t c = src->cols-1; c >= 0; c--)
    {
        setUint8Pixel(dst, c, r, getUint8Pixel(src, c, r));
    }
}
```

Performance optimization

Loop unrolling

- *Testing conditions of a loop takes instructions and hence execution time*

$$160 \times 120 = 19200$$

```
uint8_pixel_t min=UINT8_PIXEL_MAX, max=UINT8_PIXEL_MIN;
uint32_t imsize = src->rows * src->cols;
uint8_pixel_t *s = (uint8_pixel_t *)src->data;

// Scan input image for min/max values
for(uint32_t i=0; i<imsize; ++i)
{
    if(*s < min){ min = *s; }
    if(*s > max){ max = *s; }
    ++s;
}
```

Performance optimization

Loop unrolling

- *Testing conditions of a loop takes instructions and hence execution time*

$$\frac{160 \times 120}{2} = 9600$$

```
uint8_pixel_t min=UINT8_PIXEL_MAX, max=UINT8_PIXEL_MIN;
uint32_t imsize = (src->rows * src->cols) / 2;
uint8_pixel_t *s = (uint8_pixel_t *)src->data;

// Scan input image for min/max values
for(uint32_t i=0; i<imszie; ++i)
{
    if(*s < min){ min = *s; }
    if(*s > max){ max = *s; }
    ++s;

    if(*s < min){ min = *s; }
    if(*s > max){ max = *s; }
    ++s;
}
```

Performance optimization

Loop unrolling

- *Testing conditions of a loop takes instructions and hence execution time*

0

```
uint8_pixel_t min=UINT8_PIXEL_MAX, max=UINT8_PIXEL_MIN;  
  
uint8_pixel_t *s = (uint8_pixel_t *)src->data;  
  
// Scan input image for min/max values  
if(*s < min){ min = *s; }  
if(*s > max){ max = *s; }  
++s;  
  
if(*s < min){ min = *s; }  
if(*s > max){ max = *s; }  
++s;  
  
...  
  
if(*s < min){ min = *s; }  
if(*s > max){ max = *s; }  
++s;
```

Performance optimization

Perform a calculation only a single time, however...

```
// Scale the output to basic image type
for(uint32_t i=0; i<imsize; ++i)
{
    *d++ = (uint8_pixel_t)((255.0f/(max-min)) * (*s++ - min) + 0.5f);
}
```

Optimize most (-O3)
~**3520 μs**

```
// Scale the output to basic image type
float factor = 255.0f/(max-min);

for(uint32_t i=0; i<imsize; ++i)
{
    *d++ = (uint8_pixel_t)((factor) * (*s++ - min) + 0.5f);
}
```

Optimize most (-O3)
~**3520 μs**

Performance optimization

Construct a lookup table (LUT)

- All src pixels with the same value, will get the same new value after a calculation

		src						dst			
		0	1	2	3			0	1	2	3
0	0	4	4	4	4	0	0	0	0	0	
	1	5	5	5	5		85	85	85	85	
2	6	6	6	6	6	2	170	170	170	170	
3	7	7	7	7	7	3	255	255	255	255	

$$p_{dst}(x, y) = factor \cdot (p_{src}(x, y) - src_{min}) + 0.5$$

- Instead of a calculation for each pixel and writing the result to the dst the image...

Performance optimization

Construct a lookup table (LUT)

- All src pixels with the same value, will get the same new value after a calculation

		src						dst			
		0	1	2	3			0	1	2	3
0	0	4	4	4	4	0	0	0	0	0	
	1	5	5	5	5		85	85	85	85	
2	6	6	6	6	6	2	170	170	170	170	
3	7	7	7	7	7	3	255	255	255	255	

$$LUT[\text{index}] = \text{factor} \cdot (\text{index} - \text{src}_{\min}) + 0.5$$

- Construct a table (an array), where *index* are all graylevels in src

Performance optimization

Construct a lookup table (LUT)

- All src pixels with the same value, will get the same new value after a calculation

		src						dst			
		0	1	2	3			0	1	2	3
index	0	4	4	4	4	0	0	0	0	0	
	1	5	5	5	5		85	85	85	85	
2	6	6	6	6	6	1	170	170	170	170	
3	7	7	7	7	7	2	255	255	255	255	

index	0	1	2	3	4	5	6	7	8	9	255	Gray values in src
LUT	0	0	0	0	0	85	170	255	0	0	0	Gray values in dst
								...	0	0	0	

- The time-consuming calculations is thus performed only 256 times

Performance optimization

Construct a lookup table (LUT)

- All src pixels with the same value, will get the same new value after a calculation


		src						dst			
		0	1	2	3			0	1	2	3
0	0	4	4	4	4	0	0	0	0	0	
	1	5	5	5	5		85	85	85	85	
2	6	6	6	6	6	2	170	170	170	170	
3	7	7	7	7	7	3	255	255	255	255	

$$p_{dst}(x, y) = LUT[p_{src}(x, y)]$$

- And use the constructed table to assign values to dst

Performance optimization

Use for a (local) counter 32-bit variables !



- Example: count from 0 to 100
- Using a `uint8_t` is not efficient in a 32-bit microcontroller
- All registers and RAM are 32-bit.
- When using an `uint8_t`, the compiler must make sure that $255 + 1 = 0$
- This is an additional instruction (which one?) for each addition!!

```
uint8_t i = 0;  
  
while(i < 100)  
{  
    // Do work here  
    i++;  
}
```

Performance optimization

Read/write as less as possible from/to RAM!

- If possible, read/write 4 pixels in a single cycle

```
// Copy uint8_t image in chunks of four pixels
long int i = src->rows * src->cols / 4;
uint32_t *s = (uint32_t *)src->data;
uint32_t *d = (uint32_t *)dst->data;

dst->rows = src->rows;
dst->cols = src->cols;
dst->type = src->type;

while(i-- > 0)
{
    *d++ = *s++;
}
```

Performance optimization

Read/write as less as possible from/to RAM!

- If possible, read/write 4 pixels in a single cycle
- Operations are executed on registers in the CPU.
- Memory access is a time consuming (and energy intensive) operation.
- Bit-shifting a 32-bit variable takes less execution time than reading 4 bytes from memory!

```
// src is a uint8_pixel_t image
uint32_t *s = (uint32_t *)src->data;

uint32_t four_pixels = *s++;

if((uint8_pixel_t)(four_pixels >> 0) > max){max = (uint8_pixel_t)(four_pixels >> 0);}
if((uint8_pixel_t)(four_pixels >> 8) > max){max = (uint8_pixel_t)(four_pixels >> 8);}
Etc.
```

Performance optimization

Use mipmaps (image pyramids)

- Example: `scale()` (Optimize most (-O3))
 - 160x120: ~3520 µs
 - 80x60: ~880 µs
 - 40x30: ~220 µs
 - 20x15: ~60 µs
- ÷ 4
÷ 4
÷ 4




Not useful if no grayscale information should be lost

Performance optimization

Use mipmaps (image pyramids)

- Example: `threshold()` (Optimize most (-O3))
 - 160x120: ~4570 µs
 - 80x60: ~1150 µs
 - 40x30: ~290 µs
 - 20x15: ~80 µs
- ÷ 4
÷ 4
÷ 4



Very useful for finding (the location of) binary objects

Performance optimization



Use mipmaps (image pyramids)

- *Example: 160x120: ~16000 µs*

threshold()



labelTwoPass()



Performance optimization


Use mipmaps (image pyramids)

- *Example: 80x60: ~4200 µs*

threshold()



labelTwoPass()



Performance optimization

Use mipmaps (image pyramids)

- *Example: 40x30: ~1010 µs*

threshold()



labelTwoPass()



Performance optimization

Use mipmaps (image pyramids)

- *Example: 20x15:* ~280 µs

threshold()



labelTwoPass()



*An object located at (10, 10)
is located approximately at (80, 80)
in the original format*

Performance optimization

Use mipmaps (image pyramids)

- Use the *zoomFactor()* function for creating a pyramid

```
// -----
// Local image memory allocation
// -----
image_t *src = newUint8Image(EVDK5_WIDTH, EVDK5_HEIGHT);
image_t *dst = newUint8Image(EVDK5_WIDTH, EVDK5_HEIGHT);

const uint32_t zoom_factor = 2;
image_t *tmp = newUint8Image(EVDK5_WIDTH/zoom_factor, EVDK5_HEIGHT/zoom_factor);
```

Performance optimization

Use mipmaps (image pyramids)

- Use the *zoomFactor()* function for creating a pyramid

```
while(1U)
{
    // -----
    // Wait for camera image complete
    // -----
    while(smartdma_camera_image_complete == 0)
    {}

    smartdma_camera_image_complete = 0;

    // -----
    // Image processing pipeline
    // -----
    // Convert uyvy_pixel_t camera image to uint8_pixel_t image
    convertToInt8(cam, src);
```

Performance optimization

Use mipmaps (image pyramids)

- Use the *zoomFactor()* function for creating a pyramid


```
// Zoom the image
zoomFactor(src, tmp, 0, 0, src->cols, src->rows, ZOOM_OUT, zoom_factor);

// Process the zoomed image
threshold(tmp, tmp, 0, 64);
labelTwoPass(tmp, tmp, CONNECTED_EIGHT, 64);
scale(tmp, tmp);
```

Performance optimization

Use mipmaps (image pyramids)

- If needed, show the zoomed image by copying it to a QQVGA image (which is the size of the usb image)
- Add a border for nice visual effect



Performance optimization

Use mipmaps (image pyramids)

- If needed, show the zoomed image by copying it to a QQVGA image (which is the size of the usb image)
- Add a border for nice visual effect

```
// Show zoomed image
// Only one resolution can be transferred via USB, so:
// 1. Set destination image to gray background (64) for nicer visual effect
// 2. Copy smaller image to the center of the larger image for
//    nicer visual effect
memset(dst->data, 64, dst->cols * dst->rows);
int offset_c = (dst->cols / 2) - (tmp->cols / 2);
int offset_r = (dst->rows / 2) - (tmp->rows / 2);
for(int r=0; r < tmp->rows; r++)
{
    memcpy((dst->data)+((r+offset_r) * dst->cols) + offset_c, (tmp->data) +
           (r * tmp->cols), tmp->cols);
}
```

Performance optimization

Use mipmaps (image pyramids)

- Convert to USB image and set the flag

```
// Convert uint8_pixel_t image to bgr888_pixel_t image for USB
convertToBgr888(dst, usb);

// -----
// Set flag for USB interface that a new frame is available
// -----
image_available_for_usb = 1;
}
```

References

- Shore, C. (2010) Efficient C Code for ARM Devices, ARM, Downloaded November 2022 from
https://community.arm.com/cfs-file/_key/telligent-evolution-components-attachments/01-2142-00-00-00-01-26-13/ATC_2D00_152_5F00_paper_5F00_Shore.pdf
- Mukherjee, S. () Efficient C Code for ARM Devices, ARM, Downloaded November 2022 from
<https://web.archive.org/web/20170829213827/https://www.arm.com/files/pdf/AT - Better C Code for ARM Devices.pdf>
- Wikipedia contributors. (2024, August 24). Mipmap. In Wikipedia, The Free Encyclopedia. Retrieved 18:58, October 5, 2024, from
<https://en.wikipedia.org/w/index.php?title=Mipmap&oldid=1242024792>

EVD1 – Assignment



Study guide
Week 2

1 Image fundamentals – scaleFast()

Rotate 180

- The camera is mounted upside down
- We need an ultra-fast implementation for rotating the image 180 degrees



Rotate 180

- The following methods will be compared for performance
 1. Using gono functions
 2. Flipping the image in both horizontal and vertical direction
 - a) C
 - b) ARM inline assembly
 - c) ARM with C idiom
 - d) ARM 32-bit architecture optimized assembly



Rotate 180 - gonio


- Rotation is defined as:

$$x' = x_c + (x - x_c) \cdot \cos\theta - (y - y_c) \cdot \sin\theta$$

$$y' = y_c + (x - x_c) \cdot \sin\theta + (y - y_c) \cdot \cos\theta$$

where

(x_c, y_c) : the rotation origin



Rotate 180 - gonio


- Rotation is defined as:

$$x' = x_c + (x - x_c) \cdot \cos\theta - (y - y_c) \cdot \sin\theta$$

$$y' = y_c + (x - x_c) \cdot \sin\theta + (y - y_c) \cdot \cos\theta$$

Translate

Translate




Rotate 180 - gonio

- Rotation is defined as:

$$x' = x_c + (x - x_c) \cdot \cos\theta - (y - y_c) \cdot \sin\theta$$

$$y' = y_c + (x - x_c) \cdot \sin\theta + (y - y_c) \cdot \cos\theta$$


Rotate Rotate



Rotate 180 - gonio


- Rotation is defined as:

$$x' = x_c + (x - x_c) \cdot \cos\theta - (y - y_c) \cdot \sin\theta$$

$$y' = y_c + (x - x_c) \cdot \sin\theta + (y - y_c) \cdot \cos\theta$$



Translate



Rotate 180 - gonio


```
void rotate( const image_t *src, image_t *dst, const float radians,  
            const point_t center);
```

See file **EVDK_Operators\graphics_algorithms.c**

Implementation	Execution time (QQVGA and optimize most (-O3)):
Gonio	11330 us
Flip in C	
Flip in ARM inline assembly	
Flip in ARM with C idiom	
Flip in ARM architecture optimized assembly	


Rotate 180 – Flip in C

- Flip the image in both directions




Rotate 180 – Flip in C

- Flip the image in both directions



Rotate 180 – Flip in C

- Flip the image in both directions



```
uint32_t i;
uint8_pixel_t *s = (uint8_pixel_t *)img->data;
uint8_pixel_t *d = (uint8_pixel_t *)img->data +
    (img->rows * img->cols * sizeof(uint8_pixel_t)) -
    (1 * sizeof(uint8_pixel_t));


uint8_pixel_t t;

for(i = (img->rows * img->cols) / 2; i > 0; i--)
{
    t      = *s; // 1
}
```

- Uses two pointers to iterate (half) the image

Rotate 180 – Flip in C

- Flip the image in both directions



```
uint32_t i;
uint8_pixel_t *s = (uint8_pixel_t *)img->data;
uint8_pixel_t *d = (uint8_pixel_t *)img->data +
    (img->rows * img->cols * sizeof(uint8_pixel_t)) -
    (1 * sizeof(uint8_pixel_t));


uint8_pixel_t t;

for(i = (img->rows * img->cols) / 2; i > 0; i--)
{
    t     = *s; // 1
    *s++ = *d; // 2
}
```

- Uses two pointers to iterate (half) the image

Rotate 180 – Flip in C

- Flip the image in both directions



```
uint32_t i;
uint8_pixel_t *s = (uint8_pixel_t *)img->data;
uint8_pixel_t *d = (uint8_pixel_t *)img->data +
    (img->rows * img->cols * sizeof(uint8_pixel_t)) -
    (1 * sizeof(uint8_pixel_t));

uint8_pixel_t t;

for(i = (img->rows * img->cols) / 2; i > 0; i--)
{
    t      = *s; // 1
    *s++ = *d; // 2
    *d-- = t;  // 3
}
```

- Uses two pointers to iterate (half) the image

Rotate 180 – Flip in C

```
void rotate180_c(const image_t *img);
```

See file **EVDK_Operators\graphics_algorithms.c**

Implementation	Execution time (QQVGA and optimize most (-O3)):
Gonio	11330 us
Flip in C	490 us
Flip in ARM inline assembly	
Flip in ARM with C idiom	
Flip in ARM architecture optimized assembly	

Rotate 180 – Flip in ARM inline assembly




```
// Pointer to the first four pixels
register uint32_t *first_ptr = (uint32_t *)img->data;

// Pointer to the end of the data
register uint32_t *last_ptr = (uint32_t *)(img->data + (img->rows * img->cols * sizeof(uint8_pixel_t)));


// Temporary variables
register uint32_t first_pixels, last_pixels;
```

Rotate 180 – Flip in ARM inline assembly



```
while(first_ptr != last_ptr)
{
    // Read pixels
    first_pixels = *first_ptr; // 1
    last_pixels = *(--last_ptr); // 2
```


Rotate 180 – Flip in ARM inline assembly



```
// Reverse 32-bit byte order : b3 b2 b1 b0 -> b0 b1 b2 b3
__asm__ ("REV %[result], %[value]" : [result] "=r" (first_pixels) : [value] "r" (first_pixels)); // 3
__asm__ ("REV %[result], %[value]" : [result] "=r" (last_pixels) : [value] "r" (last_pixels)); // 4
```

<https://developer.arm.com/documentation/100235/0100/The-Cortex-M33-Instruction-Set/Cortex-M33-instructions?lang=en>

Rotate 180 – Flip in ARM inline assembly



```
*(first_ptr++) = last_pixels; // 5  
*last_ptr      = first_pixels; // 6
```

```
}
```

Rotate 180 – Flip in ARM inline assembly

```
void rotate180_arm( const image_t *img);
```

See file **EVDK_Operators\graphics_algorithms.c**

Implementation	Execution time (QQVGA and optimize most (-O3)):
Gonio	11330 us
Flip in C	490 us
Flip in ARM inline assembly	160 us
Flip in ARM with C idiom	
Flip in ARM architecture optimized assembly	

Rotate 180 – Flip in ARM with C idiom

Some compilers at a specific optimization level recognise this pattern in C! This is called **Idiom recognition**.

```
// Reverse 32-bit byte order : b3 b2 b1 b0 -> b0 b1 b2 b3
first_pixels = ((first_pixels >> 24) & 0x000000FF) |
                ((first_pixels >> 8) & 0x0000FF00) |
                ((first_pixels << 8) & 0x00FF0000) |
                ((first_pixels << 24) & 0xFF000000);

last_pixels = ((last_pixels >> 24) & 0x000000FF) |
                ((last_pixels >> 8) & 0x0000FF00) |
                ((last_pixels << 8) & 0x00FF0000) |
                ((last_pixels << 24) & 0xFF000000);
```

```
// Reverse 32-bit byte order : b3 b2 b1 b0 -> b0 b1 b2 b3
__asm__ ("REV %[result], %[value]" : [result] "=r" (first_pixels) : [value] "r" (first_pixels));
__asm__ ("REV %[result], %[value]" : [result] "=r" (last_pixels) : [value] "r" (last_pixels));
```

Rotate 180 – Flip in ARM with C idiom

```
void rotate180_arm( const image_t *img);
```


See file **EVDK_Operators\graphics_algorithms.c**

Implementation	Execution time (QQVGA and optimize most (-O3)):
Gonio	11330 us
Flip in C	490 us
Flip in ARM inline assembly	160 us
Flip in ARM with C idiom	160 us
Flip in ARM architecture optimized assembly	

Rotate 180 – Flip in ARM architecture optimized assembly

- Minimize the number of load and store operations
- Minimize the number of transfers between core registers


```
// Note: For using this function declare the
// following function prototype external in the
// project:
//
// extern void rotate180_cm33(const image_t *img);
//
// Example usage:
//
// rotate180_cm33(img);
```



- Core register r0 holds the pointer to the image, because there is one argument


Rotate 180 – Flip in ARM architecture optimized assembly

```
rotate180_cm33:  
  
    PUSH {r4-r10}  
  
    // Load four words from image pointer  
    // r0 = image_t.cols  
    // r1 = image_t.rows  
    // r2 = image_t.type  
    // r3 = image_t.data  
    LDMIA r0, {r0-r3}
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rotate180_cm33:  
  
    PUSH {r4-r10}  
  
    // Load four words from image pointer  
    // r0 = image_t.cols  
    // r1 = image_t.rows  
    // r2 = image_t.type  
    // r3 = image_t.data  
    LDMIA r0, {r0-r3}  
  
    // Image size = cols x rows  
    MUL r1, r0, r1
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rotate180_cm33:  
  
    PUSH {r4-r10}  
  
    // Load four words from image pointer  
    // r0 = image_t.cols  
    // r1 = image_t.rows  
    // r2 = image_t.type  
    // r3 = image_t.data  
    LDMIA r0, {r0-r3}  
  
    // Image size = cols x rows  
    MUL r1, r0, r1  
  
    // Backward pointer = size + data pointer  
    ADD r1, r1, r3
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rotate180_cm33:  
  
    PUSH {r4-r10}  
  
    // Load four words from image pointer  
    // r0 = image_t.cols  
    // r1 = image_t.rows  
    // r2 = image_t.type  
    // r3 = image_t.data  
    LDMIA r0, {r0-r3}  
  
    // Image size = cols x rows  
    MUL r1, r0, r1  
  
    // Backward pointer = size + data pointer  
    ADD r1, r1, r3  
  
    // Forward pointer = data pointer  
    MOV r0, r3
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rev_loop:  
    // Load four words from forward pointer  
    // and increment address afterwards  
    LDMIA r0, {r2-r5}
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rev_loop:  
    // Load four words from forward pointer  
    // and increment address afterwards  
    LDMIA r0, {r2-r5}  
  
    // Load four words from backward pointer,  
    // decrement address before and update  
    LDMDB r1!, {r6-r9}
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rev_loop:  
    // Load four words from forward pointer  
    // and increment address afterwards  
    LDMIA r0, {r2-r5}  
  
    // Load four words from backward pointer,  
    // decrement address before and update  
    LDMDB r1!, {r6-r9}  
  
    // Reverse the bytes in each word, as well  
    // as the words themselves  
    REV r10, r2
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rev_loop:  
    // Load four words from forward pointer  
    // and increment address afterwards  
    LDMIA r0, {r2-r5}  
  
    // Load four words from backward pointer,  
    // decrement address before and update  
    LDMDB r1!, {r6-r9}  
  
    // Reverse the bytes in each word, as well  
    // as the words themselves  
    REV r10, r2  
    REV r2, r9
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rev_loop:  
    // Load four words from forward pointer  
    // and increment address afterwards  
    LDMIA r0, {r2-r5}  
  
    // Load four words from backward pointer,  
    // decrement address before and update  
    LDMDB r1!, {r6-r9}  
  
    // Reverse the bytes in each word, as well  
    // as the words themselves  
    REV r10, r2  
    REV r2, r9  
    REV r9, r3
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rev_loop:  
    // Load four words from forward pointer  
    // and increment address afterwards  
    LDMIA r0, {r2-r5}  
  
    // Load four words from backward pointer,  
    // decrement address before and update  
    LDMDB r1!, {r6-r9}  
  
    // Reverse the bytes in each word, as well  
    // as the words themselves  
    REV r10, r2  
    REV r2, r9  
    REV r9, r3  
    REV r3, r8
```




Rotate 180 – Flip in ARM architecture optimized assembly

```
rev_loop:  
    // Load four words from forward pointer  
    // and increment address afterwards  
    LDMIA r0, {r2-r5}  
  
    // Load four words from backward pointer,  
    // decrement address before and update  
    LDMDB r1!, {r6-r9}  
  
    // Reverse the bytes in each word, as well  
    // as the words themselves  
    REV r10, r2  
    REV r2, r9  
    REV r9, r3  
    REV r3, r8  
    REV r8, r4  
    REV r4, r7  
    REV r7, r5  
    REV r5, r6
```



Rotate 180 – Flip in ARM architecture optimized assembly

```
// Write transformed words to forward  
// pointer, increment address afterwards  
// and update  
STMIA r0!, {r2-r5}  
  
// Write transformed words back to  
// backward pointer and increment address  
STMIA r1, {r7-r10}  
  
// Repeat until forward and backward  
// pointer meet  
CMP r0, r1  
BNE rev_loop  
  
POP {r4-r10}  
BX lr
```



Rotate 180 – Flip in ARM architecture optimized assembly

```
void rotate180_cm33(      const image_t *img);
```

See file **source\rotate180_cm33.s** (target projects only)

Implementation	Execution time (QQVGA and optimize most (-O3)):
Gonio	11330 us
Flip in C	490 us
Flip in ARM inline assembly	160 us
Flip in ARM with C idiom	160 us
Flip in ARM architecture optimized assembly	120 us

EVD1 – Assignment



Study guide **Week 2**

- 2 Image fundamentals – clearUInt8Image_cm33()
- 3 EXTRA Image fundamentals – convertUyvyToUInt8_cm33()

Graphics Algorithms

- Manipulate the position of pixels in an image
- Used to create images by using algorithms
- Affine transformation
- Warp

Affine transformation

- What is an affine transformation?

“An affine transformation is a function between affine spaces which preserves points, straight lines and planes”

- An affine transformation changes a pixel location. It does not affect its value.
- Characteristics
 - Lines map to lines
 - Parallel lines remain parallel
 - Ratios are preserved

Affine transformation

- Translation example:

$$x' = x + c \text{ (where } c = 50)$$

$$y' = y + f \text{ (where } f = 50)$$



Affine transformation

- Scaling example:

$$x' = ax \text{ (where } a = 2)$$

$$y' = ey \text{ (where } e = 2)$$



Affine transformation

- Conclusion: there is a function T that transforms source coordinates to destination coordinates:

$$(x', y') = T \cdot (x, y)$$

- In a matrix notation:

$$\begin{vmatrix} x' \\ y' \end{vmatrix} = T \begin{vmatrix} x \\ y \end{vmatrix}$$

- Scaling example:

$$\begin{vmatrix} ax \\ ey \end{vmatrix} = T \begin{vmatrix} x \\ y \end{vmatrix}$$

- Solving T :

$$\begin{vmatrix} ax \\ ey \end{vmatrix} = \begin{vmatrix} a & 0 \\ 0 & e \end{vmatrix} \begin{vmatrix} x \\ y \end{vmatrix}$$

Affine transformation

- Translation example:

$$\begin{vmatrix} x + c \\ y + f \end{vmatrix} = T \begin{vmatrix} x \\ y \end{vmatrix}$$

- However, there is no solution for T ...
- To solve T , we can use homogeneous coordinates instead of Cartesian coordinates!

Affine transformation

- The relation between Cartesian coordinates (x, y) and homogeneous coordinates (X, Y, Z) is defined as:

$$x = \frac{X}{Z} \text{ and } y = \frac{Y}{Z} \text{ with } Z \neq 0$$

- When $Z = 1$, we get:

$$x = X \text{ and } y = Y$$

- In a matrix notation when $Z = 1$:

$$\begin{vmatrix} X' \\ Y' \\ Z' \end{vmatrix} = T \begin{vmatrix} X \\ Y \\ Z \end{vmatrix} \equiv \begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = T \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

Affine transformation

- For translation this yields:

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} x + c \\ y + f \\ 1 \end{vmatrix} = T \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

- Now, solving T is possible:

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} x + c \\ y + f \\ 1 \end{vmatrix} = \begin{vmatrix} 1 & 0 & \mathbf{c} \\ 0 & 1 & \mathbf{f} \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

- Rewritten:

$$\begin{aligned} x' &= 1x + 0y + 1c = x + c \rightarrow c: "translation over x" \\ y' &= 0x + 1y + 1f = y + f \rightarrow f: "translation over y" \end{aligned}$$

Affine transformation

- Now back to scaling.
How can we also define scaling with homogeneous coordinates?

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} a & 0 & 0 \\ 0 & e & 0 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

- Rewritten:

$$\begin{aligned} x' &= ax + 0y + 0 = ax \\ y' &= 0x + ey + 0 = ey \end{aligned}$$

Affine transformation

- What does the following transformation matrix accomplish?

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} 1 & \mathbf{b} & 0 \\ \mathbf{d} & 1 & 0 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

- Rewritten:

$$\begin{aligned} x' &= x + by \\ y' &= dx + y \end{aligned}$$

- dst x is based on src x plus a factor b times y
- dst y is based on src y plus a factor d times x
- This operation is known as shearing

Affine transformation

- Conclusion: a transformation matrix is used to calculate the pixel location (x', y') in the destination image from the original pixel location (x, y) in the source image
- The general form of an affine transformation matrix is defined as:

$$T = \begin{vmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{vmatrix}$$


Affine transformation - examples

- Examples: identity

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

Rewritten:

$$\begin{aligned} x' &= x + 0y + 0 = x \\ y' &= 0x + y + 0 = y \end{aligned}$$



Affine transformation - examples

- Examples: scale

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} a & 0 & 0 \\ 0 & e & 0 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

where


$$a = 2$$

$$e = 2$$

Rewritten:

$$x' = 2x + 0y + 0 = 2x$$

$$y' = 0x + 2y + 0 = 2y$$



Affine transformation - examples

- Examples: Translation

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} 1 & 0 & c \\ 0 & 1 & f \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

where


$$c = 50$$

$$f = 50$$

Rewritten:

$$x' = x + 0y + 50 = x + 50$$

$$y' = 0x + y + 50 = y + 50$$



Affine transformation - examples

- Examples: Shear

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} 1 & b & 0 \\ d & 1 & 0 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

where

$$b = 0.50$$

$$d = 0.25$$

Rewritten:

$$x' = x + 0.5y + 0 = x + 0.5y$$

$$y' = 0.25x + y + 0 = 0.25x + y$$




Affine transformation - examples

- Examples: Translate and scale combined

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} 2 & 0 & 20 \\ 0 & 2 & 20 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

Rewritten:

$$\begin{aligned} x' &= 2x + 0y + 20 = 2x + 20 \\ y' &= 0x + 2y + 20 = 2y + 20 \end{aligned}$$



Affine transformation - examples

- Examples: Rotate CCW

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} \cos\theta & \sin\theta & 0 \\ -\sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

Where

$$\theta = \frac{\pi}{4} \text{ rad}$$

Rewritten:

$$\begin{aligned} x' &= x\cos\theta + y\sin\theta + 0 \\ y' &= -x\sin\theta + y\cos\theta + 0 \end{aligned}$$



Affine transformation - examples

- Examples: Rotate CW


$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = \begin{vmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

Where

$$\theta = \frac{\pi}{4} \text{ rad}$$

Rewritten:

$$\begin{aligned} x' &= x\cos\theta - y\sin\theta + 0 \\ y' &= x\sin\theta + y\cos\theta + 0 \end{aligned}$$



Affine transformation - examples

- CW or CCW rotation ??
- Notice that in images the y-axis points downward
- So, if [Wikipedia](#) (for example) says that this rotation matrix


$$\begin{vmatrix} \cos\theta & -\sin\theta & c \\ \sin\theta & \cos\theta & f \\ 0 & 0 & 1 \end{vmatrix}$$

causes a CCW rotation, in image processing we will see a CW rotation!

Affine transformation

- What happened with the affine rotated the image?


 affine - forward



	0	1	2	3	4	5	6	7	8	9	10	11
0	21	0	0	0	0	0	0	0	0	0	0	0
1	17	19	0	0	0	0	0	0	0	0	0	0
2	17	17	19	0	0	0	0	0	0	0	0	0
3	19	18	0	24	0	0	0	0	0	0	0	0
4	20	42	30	71	77	0	0	0	0	0	0	0
5	27	56	77	0	76	60	0	0	0	0	0	0
6	74	81	72	75	47	34	33	0	0	0	0	0
7	80	68	0	41	34	0	33	32	0	0	0	0
8	77	35	37	22	33	33	32	32	32	0	0	0
9	49	34	33	0	33	33	33	0	32	31	0	0
10	43	34	0	33	32	0	32	33	32	0	32	0
11	34	32	33	33	32	32	32	33	33	33	32	32
12	34	33	32	0	32	32	33	0	33	33	0	32

Affine transformation

- Solution 1: bilinear interpolation
- Solution 2: average filter
- Solution 3: backward transformation instead of forward transformation



Affine transformation

- Solution 1: bilinear interpolation
- Solution 2: average filter
- Solution 3: backward transformation

$$(x', y') = T \cdot (x, y) \quad \Rightarrow \quad (x, y) = \frac{(x', y')}{T}$$

Affine transformation

- In matrix calculation, division is performed by multiplication with the inverse matrix

$$\begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} = T \begin{vmatrix} x \\ y \\ 1 \end{vmatrix} \Rightarrow \begin{vmatrix} x' \\ y' \\ 1 \end{vmatrix} T^{-1} = \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$

Affine transformation

- Identity

$$T^{-1} = \begin{vmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{vmatrix}^{-1} = \begin{vmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{vmatrix}$$

- Scaling

$$T^{-1} = \begin{vmatrix} a & 0 & 0 \\ 0 & e & 0 \\ 0 & 0 & 1 \end{vmatrix}^{-1} = \begin{vmatrix} 1/a & 0 & 0 \\ 0 & 1/e & 0 \\ 0 & 0 & 1 \end{vmatrix}$$

wolframalpha.com

- Translation

$$T^{-1} = \begin{vmatrix} 1 & 0 & c \\ 0 & 1 & f \\ 0 & 0 & 1 \end{vmatrix}^{-1} = \begin{vmatrix} 1 & 0 & -c \\ 0 & 1 & -f \\ 0 & 0 & 1 \end{vmatrix}$$

Affine transformation

- Shear

$$T^{-1} = \begin{vmatrix} 1 & b & 0 \\ d & 1 & 0 \\ 0 & 0 & 1 \end{vmatrix}^{-1} = \begin{vmatrix} 1/(1-ab) & a/(ab-1) & 0 \\ b/(ab-1) & 1/(1-ab) & 0 \\ 0 & 0 & 1 \end{vmatrix}$$

- Rotate CW

$$T^{-1} = \begin{vmatrix} \cos\theta & -\sin\theta & c \\ \sin\theta & \cos\theta & f \\ 0 & 0 & 1 \end{vmatrix}^{-1} = \begin{vmatrix} \cos\theta & \sin\theta & (-c \cdot \cos\theta - f \cdot \sin\theta) \\ -\sin\theta & \cos\theta & (c \cdot \sin\theta - f \cdot \cos\theta) \\ 0 & 0 & 1 \end{vmatrix}$$

- Translation and scale combined

$$T^{-1} = \begin{vmatrix} a & 0 & c \\ 0 & e & f \\ 0 & 0 & 1 \end{vmatrix}^{-1} = \begin{vmatrix} 1/a & 0 & -c/a \\ 0 & 1/e & -f/e \\ 0 & 0 & 1 \end{vmatrix}$$

Affine transformation

- Backward affine transformation

affine - reverse



	0	1	2	3	4	5	6	7	8	9	10	11
0	21	19	19	0	0	0	0	0	0	0	0	0
1	17	19	19	19	0	0	0	0	0	0	0	0
2	17	17	19	24	24	0	0	0	0	0	0	0
3	17	18	30	24	61	77	0	0	0	0	0	0
4	22	42	30	71	77	77	60	0	0	0	0	0
5	66	56	77	76	76	60	35	35	0	0	0	0
6	66	81	72	75	54	34	35	33	32	0	0	0
7	80	68	68	41	34	34	33	32	32	32	0	0
8	55	61	37	33	34	33	32	32	32	32	31	0
9	55	34	33	33	33	33	33	32	32	31	31	32
10	34	34	34	33	32	32	32	33	32	32	32	32
11	33	33	33	32	32	32	32	33	32	33	32	32
12	33	33	32	32	32	32	33	33	33	33	32	32

Affine transformation - algorithm

```
void affineTransformation(    const image_t *src, image_t *dst,  
                            eTransformDirection d, float m[][3]);
```

See file **EVDK_Operators\graphics_algorithms.c**

```
// Scale example  
float a = 2.0f;  
float b = 2.0f;  
  
float m_forward[2][3] = {{ a , 0, 0},  
                         { 0 , b, 0}};  
float m_backward[2][3] = {{ 1/a, 0 , 0},  
                          { 0 , 1/b, 0}};  
  
affineTransformation(src, dst, TRANSFORM_BACKWARD, m_backward);
```

Warp

- What does the following transformation matrix accomplish?

$$\begin{vmatrix} X \\ Y \\ Z \end{vmatrix} = \begin{vmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{vmatrix} \begin{vmatrix} x \\ y \\ 1 \end{vmatrix}$$


- Rewritten:

$$\begin{aligned} X &= ax + by + c \\ Y &= dx + ey + f \\ Z &= gx + hy + 1 \end{aligned}$$

- This operation is known as geometric distortion
- Some characteristics are:
 - Lines map to lines
 - Parallel lines do not necessarily remain parallel
 - Ratios are not necessarily preserved

Warp


Example



Warp - example


From:

$$\begin{aligned}(x_0, y_0) &= (20, 0) \\(x_1, y_1) &= (60, 10) \\(x_2, y_2) &= (50, 40) \\(x_3, y_3) &= (0, 30)\end{aligned}$$



To:

$$\begin{aligned}(x'_0, y'_0) &= (40, 0) \\(x'_1, y'_1) &= (80, 20) \\(x'_2, y'_2) &= (100, 100) \\(x'_3, y'_3) &= (0, 119)\end{aligned}$$



Warp

$$\begin{aligned}X &= ax + by + c \\Y &= dx + ey + f \\Z &= gx + hy + 1\end{aligned}$$

- Given a source coordinate (x, y) , the destination coordinate (x', y') is then calculated by the functions for perspective equivalence

$$x' = \frac{X}{Z} = \frac{ax + by + c}{gx + hy + 1}$$

$$y' = \frac{Y}{Z} = \frac{dx + ey + f}{gx + hy + 1}$$

Warp

- Rewritten:

$$x' = \frac{X}{Z} = \frac{ax + by + c}{gx + hy + 1} \Rightarrow ax + by + c - gxx' - hyx' = x'$$

$$y' = \frac{Y}{Z} = \frac{dx + ey + f}{gx + hy + 1} \Rightarrow dx + ey + f - gxy' - hyy' = y'$$


With these functions for perspective equivalence, the eight transformation matrix coefficients $a \dots h$ can be calculated if **four** source coordinates (x, y) and **four** destination coordinates (x', y') are known.

(Because this yields 8 equations in the 8 unknowns $a \dots h$)

Warp - example


From:

$$\begin{aligned}(x_0, y_0) &= (20, 0) \\(x_1, y_1) &= (60, 10) \\(x_2, y_2) &= (50, 40) \\(x_3, y_3) &= (0, 30)\end{aligned}$$



To:

$$\begin{aligned}(x'_0, y'_0) &= (40, 0) \\(x'_1, y'_1) &= (80, 20) \\(x'_2, y'_2) &= (100, 100) \\(x'_3, y'_3) &= (0, 119)\end{aligned}$$



Warp

Let's define these 8 coordinates as follows:

- (x_k, y_k) are the mapping coordinates in the source image
- (x'_k, y'_k) are the mapping coordinates in the destination image where $k = 0, 1, 2, 3$

And substitute these in the functions for perspective equivalence:

$$ax_k + by_k + c - gx_kx'_k - hy_kx'_k = x'_k \quad | \quad dx_k + ey_k + f - gx_ky'_k - hy_ky'_k = y'_k$$

Then the result is the following 8 equations:

$$\begin{array}{l|l} ax_0 + by_0 + c - gx_0x'_0 - hy_0x'_0 = x'_0 & dx_0 + ey_0 + f - gx_0y'_0 - hy_0y'_0 = y'_0 \\ ax_1 + by_1 + c - gx_1x'_1 - hy_1x'_1 = x'_1 & dx_1 + ey_1 + f - gx_1y'_1 - hy_1y'_1 = y'_1 \\ ax_2 + by_2 + c - gx_2x'_2 - hy_2x'_2 = x'_2 & dx_2 + ey_2 + f - gx_2y'_2 - hy_2y'_2 = y'_2 \\ ax_3 + by_3 + c - gx_3x'_3 - hy_3x'_3 = x'_3 & dx_3 + ey_3 + f - gx_3y'_3 - hy_3y'_3 = y'_3 \end{array}$$

Warp

$$\begin{aligned} ax_0 + by_0 + c - gx_0x'_0 - hy_0x'_0 &= x'_0 \\ ax_1 + by_1 + c - gx_1x'_1 - hy_1x'_1 &= x'_1 \\ ax_2 + by_2 + c - gx_2x'_2 - hy_2x'_2 &= x'_2 \\ ax_3 + by_3 + c - gx_3x'_3 - hy_3x'_3 &= x'_3 \end{aligned}$$

$$\begin{aligned} dx_0 + ey_0 + f - gx_0y'_0 - hy_0y'_0 &= y'_0 \\ dx_1 + ey_1 + f - gx_1y'_1 - hy_1y'_1 &= y'_1 \\ dx_2 + ey_2 + f - gx_2y'_2 - hy_2y'_2 &= y'_2 \\ dx_3 + ey_3 + f - gx_3y'_3 - hy_3y'_3 &= y'_3 \end{aligned}$$

Rewriting these eight equations in a 8x8 system isolates the coefficients:

*Notice that
these are all
x and y values*

$$\left(\begin{array}{ccccccccc} x_0 & y_0 & 1 & 0 & 0 & 0 & -x_0x'_0 & -y_0x'_0 \\ x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x'_1 & -y_1x'_1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 & -x_2x'_2 & -y_2x'_2 \\ x_3 & y_3 & 1 & 0 & 0 & 0 & -x_3x'_3 & -y_3x'_3 \\ 0 & 0 & 0 & x_0 & y_0 & 1 & -x_0y'_0 & -y_0y'_0 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -x_1y'_1 & -y_1y'_1 \\ 0 & 0 & 0 & x_2 & y_2 & 1 & -x_2y'_2 & -y_2y'_2 \\ 0 & 0 & 0 & x_3 & y_3 & 1 & -x_3y'_3 & -y_3y'_3 \end{array} \right) \begin{pmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \end{pmatrix} = \begin{pmatrix} x'_0 \\ x'_1 \\ x'_2 \\ x'_3 \\ y'_0 \\ y'_1 \\ y'_2 \\ y'_3 \end{pmatrix}$$


Warp

This 8x8 system can be solved for the coefficients $a \dots h$ using Gaussian elimination (amongst other methods).

For example, with a C function for solving this systems of linear equations, provided by:


Systems of Linear Equations. (n.d.). In *Mathematics Source Library C & ASM*. Retrieved June 6, 2020, from
<http://www.mymathlib.com/matrices/linearsystems/>

Warp




Warp

Use the unit square




Warp

Use the unit square




Warp

Use the unit square




Warp

Use the unit square



Warp


Use the unit square



$$\begin{pmatrix} x_0 & y_0 & 1 & 0 & 0 & 0 & -x_0x'_0 & -y_0x'_0 \\ x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x'_1 & -y_1x'_1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 & -x_2x'_2 & -y_2x'_2 \\ x_3 & y_3 & 1 & 0 & 0 & 0 & -x_3x'_3 & -y_3x'_3 \\ 0 & 0 & 0 & x_0 & y_0 & 1 & -x_0y'_0 & -y_0y'_0 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -x_1y'_1 & -y_1y'_1 \\ 0 & 0 & 0 & x_2 & y_2 & 1 & -x_2y'_2 & -y_2y'_2 \\ 0 & 0 & 0 & x_3 & y_3 & 1 & -x_3y'_3 & -y_3y'_3 \end{pmatrix} \begin{pmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \end{pmatrix} = \begin{pmatrix} x'_0 \\ x'_1 \\ x'_2 \\ x'_3 \\ y'_0 \\ y'_1 \\ y'_2 \\ y'_3 \end{pmatrix}$$

Warp

Use the unit square




$$\begin{pmatrix} 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & -x'_1 \\ 1 & 1 & 1 & 0 & 0 & 0 & -x'_2 \\ 0 & 1 & 1 & 0 & 0 & 0 & -x'_3 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 & 1 & -y'_1 \\ 0 & 0 & 0 & 1 & 1 & 1 & -y'_2 \\ 0 & 0 & 0 & 0 & 1 & 1 & -y'_3 \end{pmatrix} \begin{pmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \end{pmatrix} = \begin{pmatrix} x'_0 \\ x'_1 \\ x'_2 \\ x'_3 \\ y'_0 \\ y'_1 \\ y'_2 \\ y'_3 \end{pmatrix}$$

$$x'_0 = c$$

Warp

Use the unit square



$$x'_0 = c$$

$$x'_1 = a + c - gx'_1$$

$$x'_2 = a + b + c - gx'_2 - hx'_2$$

$$x'_3 = b + c - hx'_3$$

$$y'_0 = f$$


$$y'_1 = d + f - gy'_1$$

$$y'_2 = d + e + f - gy'_2 - hy'_2$$

$$y'_3 = e + f - hy'_3$$

Warp

Use the unit square



$$a = x'_1 - x'_0 + gx'_1$$

$$b = x'_3 - x'_0 + hx'_3$$

$$c = x'_0$$

$$d = y'_1 - y'_0 + gy'_1$$

$$e = y'_3 - y'_0 + hy'_3$$


$$f = y'_0$$

$$g = \frac{(y'_0 - y'_1 + y'_2 - y'_3)(x'_3 - x'_2) + (x'_0 - x'_1 + x'_2 - x'_3)(y'_2 - y'_3)}{(y'_1 - y'_2)(x'_3 - x'_2) - (y'_3 - y'_2)(x'_1 - x'_2)}$$

$$h = \frac{(y'_0 - y'_1 + y'_2 - y'_3)(x'_2 - x'_1) + (x'_0 - x'_1 + x'_2 - x'_3)(y'_1 - y'_2)}{(y'_1 - y'_2)(x'_3 - x'_2) - (y'_3 - y'_2)(x'_1 - x'_2)}$$

Warp


Use the unit square



$$B = \begin{vmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{vmatrix}$$

Warp

Use the unit square



$$a = x_1 - x_0 + gx_1$$

$$b = x_3 - x_0 + hx_3$$

$$c = x_0$$

$$d = y_1 - y_0 + gy_1$$

$$e = y_3 - y_0 + hy_3$$


$$f = y_0$$

$$g = \frac{(y_0 - y_1 + y_2 - y_3)(x_3 - x_2) + (x_0 - x_1 + x_2 - x_3)(y_2 - y_3)}{(y_1 - y_2)(x_3 - x_2) - (y_3 - y_2)(x_1 - x_2)}$$

$$h = \frac{(y_0 - y_1 + y_2 - y_3)(x_2 - x_1) + (x_0 - x_1 + x_2 - x_3)(y_1 - y_2)}{(y_1 - y_2)(x_3 - x_2) - (y_3 - y_2)(x_1 - x_2)}$$

Warp


Use the unit square



$$A = \begin{vmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{vmatrix}$$

Warp

Use the unit square




$$A^{-1} = \begin{vmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{vmatrix}^{-1}$$

wolframalfa.com

Warp

Use the unit square




$$A^{-1} = \begin{vmatrix} e - fh & ch - b & bf - ce \\ fg - d & a - cg & cd - af \\ dh - eg & bg - ah & ae - bd \end{vmatrix} \cdot \frac{1}{ae - afh - bd + bfg + cdh - ceg}$$

wolframalfa.com

Warp


Use the unit square for forward transformation



$$T = B \times A^{-1}$$

Warp

Use the unit square for backward transformation



$$T^{-1} = A \times B^{-1}$$

Warp

And finally, for each pixel in the destination image:

- Given a destination coordinate (x, y) , the source coordinate (x', y') is then calculated by the functions for perspective equivalence where the coefficients are given by the inverse matrix T^{-1}




$$x' = \frac{X}{Z} = \frac{ax + by + c}{gx + hy + 1}$$

$$y' = \frac{Y}{Z} = \frac{dx + ey + f}{gx + hy + 1}$$

Warp – algorithm

```
void warpPerspective(    const image_t *src, image_t *dst,  
                        const point_t *from, const point_t *to,  
                        eTransformDirection d);
```

See file **EVDK_Operators\graphics_algorithms.c**



The function additionally creates a temporary mask image to set only pixels that are inside the polygon

Warp – algorithm

```
void warpPerspective(    const image_t *src, image_t *dst,  
                        const point_t *from, const point_t *to,  
                        eTransformDirection d);
```


See file **EVDK_Operators\graphics_algorithms.c**

```
const point_t from[4] =  
{  
    {20, 0},{60, 10},{50, 40},{0, 30}  
};  
  
const point_t to[4] =  
{  
    {40, 0},{80, 20},{100, 100},{0, 119}  
};  
  
warpPerspective(src, dst, from, to, TRANSFORM_BACKWARD);
```

Warp – algorithm

```
void warpPerspectiveFast(  
    const image_t *src, image_t *dst,  
    const point_t *from,  
    eTransformDirection d);
```

See file **EVDK_Operators\graphics_algorithms.c**



The “to” points are omitted in the function prototype, so it always warps to the dst full range.

Warp – algorithm

```
void warpPerspectiveFast(    const image_t *src, image_t *dst,  
                            const point_t *from,  
                            eTransformDirection d);
```

See file **EVDK_Operators\graphics_algorithms.c**

```
const point_t from[4] =  
{  
    {20, 0},{60, 10},{50, 40},{0, 30}  
};  
  
warpPerspectiveFast(src, dst, from, TRANSFORM_BACKWARD);
```

EVD1 – Assignment



Study guide
Week 2

4 Graphics algorithms – affineTransformation()

References

- Myler, H. R., & Weeks, A. R. (2009). *The pocket handbook of image processing algorithms in C*. Prentice Hall Press.
- Affine transformation. (n.d.). In Wikipedia. Retrieved December 12, 2019, from
https://en.wikipedia.org/wiki/Affine_transformation
- Heckbert, P. S. (1989). *Fundamentals of texture mapping and image warping*. p. 9-29
- Systems of Linear Equations. (n.d.). In *Mathematics Source Library C & ASM*. Retrieved June 6, 2020, from
<http://www.mymathlib.com/matrices/linearsystems/>

References

- Wikipedia contributors. (2024, May 3). Rotation (mathematics). In *Wikipedia, The Free Encyclopedia*. Retrieved 08:04, October 7, 2024, from [https://en.wikipedia.org/w/index.php?title=Rotation_\(mathematics\)&oldid=1222018142](https://en.wikipedia.org/w/index.php?title=Rotation_(mathematics)&oldid=1222018142)