

HAN AEA - Embedded Vision & Machine Learning

EVD1 - Week 3




Image Fundamentals
Nonlinear Filters
Spatial Filters

By Hugo Arends

Image Fundamentals

- Functions for creating and deleting images
- Functions for converting images
- Functions for reading and writing pixels
- Basic image processing operators

- Discrete convolution
- Discrete correlation

Discrete convolution

- Results in a filtering operation
- Applies a filter mask to an image by convolving the filter mask with the original image
- Two-dimensional discrete convolution is defined as

$$p_{dst}(x, y) = \sum_{i=-n/2}^{i=n/2} \sum_{j=-m/2}^{j=m/2} p_{src}(x - i, y - j) \cdot p_{mask}(i + n/2, j + m/2)$$

where

$m \times n$: mask size

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)

$p_{src}(x - i, y - j)$: a pixel value in the src image within the mask

$p_{mask}(i + n/2, j + m/2)$: a pixel value in the mask image

Discrete convolution

- Results in a filtering operation
- Applies a filter mask to an image by convolving the filter mask with the original image
- Two-dimensional discrete convolution is defined as


$$p_{dst}(x, y) = \sum_{i=-n/2}^{i=n/2} \sum_{j=-m/2}^{j=m/2} p_{src}(x - i, y - j) \cdot p_{mask}(i + n/2, j + m/2)$$

All pixels in the
 $m \times n$
neighbourhood
of the src pixel at
(x, y)

Discrete convolution

- Results in a filtering operation
- Applies a filter mask to an image by convolving the filter mask with the original image
- Two-dimensional discrete convolution is defined as

$$p_{dst}(x, y) = \sum_{i=-n/2}^{i=n/2} \sum_{j=-m/2}^{j=m/2} p_{src}(x - i, y - j) \cdot p_{mask}(i + n/2, j + m/2)$$



All pixels in the $m \times n$
mask

Discrete convolution

- The corresponding pixels in the src image and mask are **flipped** in both horizontal and vertical direction.
- If the mask is not symmetrical in both horizontal and vertical direction, the mask should be flipped before performing a convolution.

Mask

a	b	c
d	e	f
g	h	i

Flipped
mask

i	h	g
f	e	d
c	b	a

No need to
flip

0	-1	0
-1	5	-1
0	-1	0



Must be
flipped

0	0	1
0	1	0
0	1	0

Discrete convolution - example

Identity



0	0	0
0	1	0
0	0	0



Discrete convolution - example

Edge



-1	-1	-1
-1	8	-1
-1	-1	-1



Discrete convolution - example

Sharpen

0	-1	0
-1	5	-1
0	-1	0



Discrete convolution - implementation

```
void convolve(    const image_t *src, image_t *dst,  
                  const image_t *msk);
```

See file **EVDK_Operators\image_fundamentals.c**

Discrete convolution - implementation

```
void convolve(const image_t *src, image_t *dst, const image_t *msk)
{
    // Loop all pixels
    for(int32_t y=0; y<src->rows; y++)
    {
        for(int32_t x=0; x<src->cols; x++)
        {
```

Discrete convolution - implementation

```
int32_t val = 0;
int32_t dr = (msk->rows/2);
int32_t dc = (msk->cols/2);

// Apply the kernel only for pixels within the image
for(int32_t j=-dr; j<=dr; j++)
{
    for(int32_t i=-dc; i<=dc; i++)
    {
        if((x-i) >= 0 &&
           (y-j) >= 0 &&
           (x-i) < src->cols &&
           (y-j) < src->rows)
        {
            val += getInt16Pixel(src,x-i,y-j) * getInt16Pixel(msk,i+dc,j+dr);
        }
    }
}
```

Discrete convolution - implementation

```
// Clip the result
if(val>INT16_PIXEL_MAX)
    val=INT16_PIXEL_MAX;

if(val<INT16_PIXEL_MIN)
    val=INT16_PIXEL_MIN;

// Store the result
setInt16Pixel(dst, x, y, val);
}

}
```

Discrete convolution - implementation

```
// Clip the result
if(val>INT16_PIXEL_MAX)
    val=INT16_PIXEL_MAX;

if(val<INT16_PIXEL_MIN)
    val=INT16_PIXEL_MIN;

// Store the result
setInt16Pixel(dst, x, y, val);
}
}
```

Optimize most (-O3)
UVC connected
~35 ms

Discrete convolution – improve performance

- Assume a 3x3 mask and ignore border pixels

```
// Apply the kernel only for pixels within the image
for(int32_t j=-dr; j<=dr; j++)
{
    for(int32_t i=-dc; i<=dc; i++)
    {
        if((x-i) >= 0 &&
           (y-j) >= 0 &&
           (x-i) < src->cols &&
           (y-j) < src->rows)
        {
            val += getInt16Pixel(src,x-i,y-j) * getInt16Pixel(msk,i+dc,j+dr);
        }
    }
}
```

No loops needed with variable
delta r and delta c (loop unrolling)

Discrete convolution – improve performance

- Assume a 3x3 mask and ignore border pixels

```
// Apply the kernel only for pixels within the image
for(int32_t j=-dr; j<=dr; j++)
{
    for(int32_t i=-dc; i<=dc; i++)
    {
        if((x-i) >= 0 &&
           (y-j) >= 0 &&
           (x-i) <  src->cols &&
           (y-j) <  src->rows)
        {
            val += getInt16Pixel(src,x-i,y-j) * getInt16Pixel(msk,i+dc,j+dr);
        }
    }
}
```

No loops needed with variable delta r and delta c (loop unrolling)

No need to check image boundaries for every pixel

Discrete convolution – improve performance

- Assume a 3x3 mask and ignore border pixels
- Avoid using getters and setters


```
// Apply the kernel only for pixels within the image
for(int32_t j=-dr; j<=dr; j++)
{
    for(int32_t i=-dc; i<=dc; i++)
    {
        if((x-i) >= 0 &&
           (y-j) >= 0 &&
           (x-i) < src->cols &&
           (y-j) < src->rows)
        {
            val += getInt16Pixel(src,x-i,y-j) * getInt16Pixel(msk,i+dc,j+dr);
        }
    }
}
```

No loops needed with variable delta r and delta c (loop unrolling)

No need to check image boundaries for every pixel

Use pointers instead

EVD1 – Assignment



Study guide
Week 3

2 Image fundamentals – convolveFast()

Discrete correlation

- Compares two images mathematically
- The result is a two-dimensional expression of equivalence
- The mask image is often referred to as the template
- The correlation is then called template matching
- Two-dimensional discrete correlation is defined as

$$p_{dst}(x, y) = \sum_{i=-n/2}^{i=n/2} \sum_{j=-m/2}^{j=m/2} p_{src}(x + i, y + j) \cdot p_{mask}(i + n/2, j + m/2)$$

where

$m \times n$: mask size


$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)

$p_{src}(x + i, y + j)$: a pixel value in the src image within the mask

$p_{mask}(i + n/2, j + m/2)$: a pixel value in the mask image


Discrete correlation – example

Binary template matching



Discrete correlation – example

Binary template matching



Discrete correlation – implementation

```
void correlate(    const image_t *src, image_t *dst,  
                  const image_t *msk);
```

See file **EVDK_Operators\image_fundamentals.c**

Discrete correlation – implementation

```
val += getInt16Pixel(src,x+i,y+j) * getInt16Pixel(msk,i+dc,j+dr);
```

Optimize most (-O3)
UVC connected
25x25 mask
> 1000 ms

Discrete correlation – improve performance

- For template matching it does not make sense to assume a 3x3 mask, because the size of a mask is at least tens by tens pixels
- Although the border pixels can be skipped, if the mask size increases, so are the number of skipped pixels
- In other words, a template matcher with a reasonable sized mask takes too long to execute on the microcontroller

Nonlinear filters

- Operate on an image by computing a given nonlinear function over a local window
- The local window can vary in size, is often a square
- Replace one specified pixel within the local window with the computed value, often the centre pixel (hence window size is an odd value)
- Are not solely used for filtering, e.g. binary erosion, binary dilation, and edge detection

Image

window size (n) = 3

m_1	m_2	m_3			
m_4	m_5	m_6			
m_7	m_8	m_9			

Nonlinear filters

- Operate on an image by computing a given nonlinear function over a local window
- The local window can vary in size, is often a square
- Replace one specified pixel within the local window with the computed value, often the centre pixel (hence window size is an odd value)
- Are not solely used for filtering, e.g. binary erosion, binary dilation, and edge detection

Image

window size (n) = 3

	m_1	m_2	m_3		
	m_4	m_5	m_6		
	m_7	m_8	m_9		

Nonlinear filters

- Operate on an image by computing a given nonlinear function over a local window
- The local window can vary in size, is often a square
- Replace one specified pixel within the local window with the computed value, often the centre pixel (hence window size is an odd value)
- Are not solely used for filtering, e.g. binary erosion, binary dilation, and edge detection

window size (n) = 3

		Image			
		m_1	m_2	m_3	
		m_4	m_5	m_6	
		m_7	m_8	m_9	

Nonlinear filters

- Operate on an image by computing a given nonlinear function over a local window
- The local window can vary in size, is often a square
- Replace one specified pixel within the local window with the computed value, often the centre pixel (hence window size is an odd value)
- Are not solely used for filtering, e.g. binary erosion, binary dilation, and edge detection

window size (n) = 3

			Image		
			m_1	m_2	m_3
			m_4	m_5	m_6
			m_7	m_8	m_9

Nonlinear filters

- Operate on an image by computing a given nonlinear function over a local window
- The local window can vary in size, is often a square
- Replace one specified pixel within the local window with the computed value, often the centre pixel (hence window size is an odd value)
- Are not solely used for filtering, e.g. binary erosion, binary dilation, and edge detection

Image

window size (n) = 3

m_1	m_2	m_3			
m_4	m_5	m_6			
m_7	m_8	m_9			

The diagram shows a 6x6 grid representing an image. A 3x3 local window is highlighted with a blue border, containing pixels m_1 , m_2 , m_3 in the top row, m_4 , m_5 , m_6 in the middle row, and m_7 , m_8 , m_9 in the bottom row. Pixel m_5 is highlighted with a green border, indicating it is the central pixel of the window.

Nonlinear filters

- Operate on an image by computing a given nonlinear function over a local window
- The local window can vary in size, is often a square
- Replace one specified pixel within the local window with the computed value, often the centre pixel (hence window size is an odd value)
- Are not solely used for filtering, e.g. binary erosion, binary dilation, and edge detection

Image

	m_1	m_2	m_3		
	m_4	m_5	m_6		
	m_7	m_8	m_9		

window size (n) = 3

The diagram shows a 3x3 grid of pixels labeled m_1 through m_9 . A blue border highlights a 3x3 window centered on pixel m_5 . Pixel m_5 is highlighted with a green border.

Nonlinear filters – implementation

Implementation is very similar to convolution

Handles border pixels by just taking the valid pixels into account in the filter specific operation

Nonlinear filters - implementation

```
void mean(const image_t *src, image_t *dst, const uint8_t n)
{
    // Loop all pixels
    for(int32_t y=0; y<src->rows; y++)
    {
        for(int32_t x=0; x<src->cols; x++)
        {
```

Nonlinear filters - implementation

```
// Initialize filter specific variables
int32_t cnt = 0;
int32_t sum = 0;

// Apply the kernel only for pixels within the image
for(int32_t j=-n/2; j<=n/2; j++)
{
    for(int32_t i=-n/2; i<=n/2; i++)
    {
        if((x+i) >= 0 &&
           (y+j) >= 0 &&
           (x+i) < src->cols &&
           (y+j) < src->rows)
        {
            // Count the number of pixels in the calculation
            cnt++;

            // Get pixel and perform filter specific calculation
            sum += getUInt8Pixel(src,x+i,y+j);
        }
    }
}
```

Nonlinear filters - implementation

```
        // Calculate and store the result
        setUint8Pixel(dst,x,y,(uint8_pixel_t)((float)sum/(float)cnt + 0.5f));
    }
}
```

Nonlinear filters - arithmetic mean

- Calculates the arithmetic mean of the pixels within the window
- The arithmetic mean is defined as

$$p_{dst}(x, y) = \frac{1}{n^2} \sum_{i=-n/2}^{i=n/2} \sum_{j=-n/2}^{j=n/2} p_{src}(x + i, y + j)$$

where

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)

$p_{src}(x + i, y + j)$: a pixel value in the src image within the window

n : the window size

Nonlinear filters - arithmetic mean

```
void mean( const image_t *src, image_t *dst, const uint8_t n);
```

See file **EVDK_Operators\nonlinear_filters.c**

Nonlinear filters - arithmetic mean

```
void mean( const image_t *src, image_t *dst, const uint8_t n);
```

See file **EVDK_Operators\nonlinear_filters.c**

```
// -----
// Image processing pipeline
// -----  
  
convertBgr888ToInt8(img, src);  
  
mean(src, dst, 3);
```

Nonlinear filters - arithmetic mean

```
void mean( const image_t *src, image_t *dst, const uint8_t n);
```

See file **EVDK_Operators\nonlinear_filters.c**

Optimize most (-O3)


UVC connected

3x3 mask: **56 ms**



5x5 mask: **136 ms**

11x11 mask: **594 ms**


Nonlinear filters - arithmetic mean examples






*Original image,
without
(additional)
noise*



*Gaussian noise
 $m = 0$
 $\sigma = 800$*



*Salt and
pepper noise
 $p = 0.05$*



*Uniform noise
 $a = 10$
 $b = 60$*

Nonlinear filters

An overview of all nonlinear filters

```
void harmonic(    const image_t *src, image_t *dst, const uint8_t n);
void maximum(    const image_t *src, image_t *dst, const uint8_t n);
void mean(        const image_t *src, image_t *dst, const uint8_t n);
void midpoint(   const image_t *src, image_t *dst, const uint8_t n);
void minimum(   const image_t *src, image_t *dst, const uint8_t n);
void range(       const image_t *src, image_t *dst, const uint8_t n);
```

Nonlinear filters – harmonic mean

- Is better at removing Gaussian type noise and preserves edges
- Removes positive outliers
- Harmonic mean filter is defined as

$$p_{dst}(x, y) = \begin{cases} 0 & \text{If any } p_{src}(x + i, y + j) = 0 \\ m & \text{otherwise} \\ \frac{1}{\sum_{i=-n/2}^{n/2} \sum_{j=-n/2}^{n/2} p_{src}(x + i, y + j)} & \end{cases}$$

where

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)
 $p_{src}(x + i, y + j)$: a pixel value in the src image within the window
 m : the number of pixels included in the summation calculation

Nonlinear filters – median

- Removes long tailed noise and salt and pepper type noise
- Has minimum blurring effect and preserves spatial details
- Can remove outlier noise from images that contain less than 50% of its pixels as outliers
- The median is defined as

$$p_{dst}(x, y) = \begin{cases} X\left[\frac{n}{2}\right] & \text{if } n \text{ is odd} \\ \frac{X\left[\frac{n-1}{2}\right] + X\left[\frac{n}{2}\right]}{2} & \text{if } n \text{ is even} \end{cases}$$

where

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)
 n : the window size

X : sorted list of values from the source image in the window

$$[window(x + i, y + j)]_{i=-n/2}^{i=n/2}_{j=-n/2}^{j=n/2}$$

Nonlinear filters – minimum

- Outputs the local minimum
- Is used to remove positive outlier noise
- Minimum filter is defined as

$$p_{dst}(x, y) = \min_{i=-n/2}^{i=n/2} \min_{j=-n/2}^{j=n/2} [\text{window}(x + i, y + j)]$$

where

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)

$\text{window}(x + i, y + j)$: all values in the src image within the window

n : the window size

Nonlinear filters – maximum

- Outputs the local maximum
- Is used to remove negative outlier noise
- Maximum filter is defined as

$$p_{dst}(x, y) = \max[\text{window}(x + i, y + j)]_{i=-n/2}^{i=n/2}_{j=-n/2}^{j=n/2}$$

where

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)

$\text{window}(x + i, y + j)$: all values in the src image within the window

n : the window size

Nonlinear filters – midpoint

- Outputs the average of the local minimum and maximum
- Used to remove short tailed noise, such as Gaussian and uniform type noise
- Midpoint filter is defined as

$$p_{dst}(x, y) = \frac{\min[\text{window}(x + i, y + j)]_{i=-n/2}^{i=n/2}{}_{j=-n/2}^{j=n/2} + \max[\text{window}(x + i, y + j)]_{i=-n/2}^{i=n/2}{}_{j=-n/2}^{j=n/2}}{2}$$

where

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)

$\text{window}(x + i, y + j)$: all values in the src image within the window

n : the window size

Nonlinear filters – range

- Outputs the difference between the local maximum and minimum
- Range filter is defined as

$$p_{dst}(x, y) = \max[\text{window}(x + i, y + j)]_{i=-n/2}^{i=n/2} - \min[\text{window}(x + i, y + j)]_{j=-n/2}^{j=n/2}$$


where

$p_{dst}(x, y)$: the calculated result pixel in the destination image at (x, y)

$\text{window}(x + i, y + j)$: all values in the src image within the window

n : the window size

EVD1 – Assignment



Study guide
Week 3

3 Nonlinear filters – meanFast()
4 Nonlinear filters – EXTRA


Spatial Filters

- Are basically discrete convolution filters
- The filter is known as a spatial mask
- Copy an image pixel-by-pixel while allowing for the effects of the pixel values in the local area (mask)
- Filter operations include multiplication, addition, and shifting
- Convolution is a time-consuming operation, so a 3x3 mask is typically used
- Typically uses odd masks and the centre of the mask for the pixel that is to be replaced, but this is by no means a requirement
- It can be implemented in parallel hardware for extremely fast execution
- Gaussian filters
- Laplacian filters
- Sobel filter

Gaussian


- Gaussian filters are masks formed from a two-dimensional Gaussian distribution
- Removes high frequency noise, but causes blurring

Gaussian - example



3x3

1	2	1
2	4	2
1	2	1



5x5

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

Gaussian - algorithm

```
void gaussianFilter_3x3( const image_t *src, image_t *dst);  
void gaussianFilter_5x5( const image_t *src, image_t *dst);
```

See file **EVDK_Operators\spatial_filters.c**

Gaussian - algorithm

```
void gaussianFilter_3x3( const image_t *src, image_t *dst);  
void gaussianFilter_5x5( const image_t *src, image_t *dst);
```

See file **EVDK_Operators\spatial_filters.c**

```
// PC app  
  
// Create additional int16_pixel_t images, because calculations with such  
// masks will not be in the uint8_pixel_t range  
image_t *src_int16 = newInt16Image(IMG_WIDTH, IMG_HEIGHT);  
image_t *dst_int16 = newInt16Image(IMG_WIDTH, IMG_HEIGHT);  
  
cv::Mat cv_src_int16(IMG_HEIGHT, IMG_WIDTH, CV_16UC1, src_int16->data);  
cv::Mat cv_dst_int16(IMG_HEIGHT, IMG_WIDTH, CV_16UC1, dst_int16->data);
```

Gaussian - algorithm

```
void gaussianFilter_3x3( const image_t *src, image_t *dst);  
void gaussianFilter_5x5( const image_t *src, image_t *dst);
```


See file **EVDK_Operators\spatial_filters.c**

```
// -----  
// Image processing pipeline  
// -----  
// Convert input image  
convertBgr888ToInt16(img, src);  
convertBgr888ToInt16(img, src_int16);  
  
// Filter  
gaussianFilter_3x3(src_int16, dst_int16);  
  
// Scale both images to uint8_pixel_t for convenient visualisation  
scale(src, src);  
scaleInt16ToUint8(dst_int16, dst);
```


Laplacian

- Gives the second derivative in two directions
- Enhances changes and only changes


Edge




First derivative



Second derivative



Laplacian - example



3x3

0	-1	0
-1	4	-1
0	-1	0



5x5

-1	-1	-1	-1	-1
-1	-1	-1	-1	-1
-1	-1	24	-1	-1
-1	-1	-1	-1	-1
-1	-1	-1	-1	-1

Laplacian - algorithm

```
void laplacian_3x3( const image_t *src, image_t *dst);  
void laplacian_5x5( const image_t *src, image_t *dst);
```

See file **EVDK_Operators\spatial_filters.c**

Sobel

- Edge detection algorithm with two results:
 1. Edge magnitude
 2. Edge direction
- Sobel magnitude and direction are defined as

$$M_{sobel}(x, y) = |G_H(x, y)| + |G_V(x, y)| \quad \varphi_{sobel}(x, y) = \tan^{-1}\left(\frac{G_V(x, y)}{G_H(x, y)}\right)$$

where

$M_{sobel}(x, y)$: the calculated magnitude in the destination image at (x, y)

$\varphi_{sobel}(x, y)$: the calculated direction in the destination image at (x, y)

$G_H(x, y)$: the pixel at (x, y) in the horizontal enhanced image

$G_V(x, y)$: the pixel at (x, y) in the vertical enhanced image

Sobel

- The horizontal enhanced image G_H is obtained by a convolution with mask

3x3


-1	-2	-1
0	0	0
1	2	1

- The vertical enhanced image G_V is obtained by a convolution with mask


3x3

-1	0	1
-2	0	2
-1	0	1

Sobel - example



Magnitude



Direction
between
 $-\frac{\pi}{2}$ and $\frac{\pi}{2}$

Sobel - algorithm

```
void sobel( const image_t *src, image_t *mag, image_t *dir);
```

See file **EVDK_Operators\spatial_filters.c**

The dir image can be omitted (NULL). If not omitted, it must be of image type **IMGTYPE_FLOAT**.

Sobel - algorithm

```
void sobel( const image_t *src, image_t *mag, image_t *dir);
```

See file **EVDK_Operators\spatial_filters.c**

The dir image can be omitted (NULL). If not omitted, it must be of image type **IMGTYPE_FLOAT**.

```
// EVDK example

// Create additional int16_pixel_t images, because calculations with such
// masks will not be in the uint8_pixel_t range
image_t *src_int16 = newInt16Image(IMG_WIDTH, IMG_HEIGHT);
image_t *dst_int16 = newInt16Image(IMG_WIDTH, IMG_HEIGHT);
```

Sobel - algorithm

```
void sobel( const image_t *src, image_t *mag, image_t *dir);
```

See file **EVDK_Operators\spatial_filters.c**

The dir image can be omitted (NULL). If not omitted, it must be of image type **IMGTYPE_FLOAT**.

```
// -----
// Image processing pipeline
// -----
// Convert uyvy_pixel_t camera image to int16_pixel_t image
convertUyvyToInt16(cam, src_int16);
```

Sobel - algorithm

```
void sobel( const image_t *src, image_t *mag, image_t *dir);
```

See file **EVDK_Operators\spatial_filters.c**

The dir image can be omitted (NULL). If not omitted, it must be of image type **IMGTYPE_FLOAT**.

```
// Copy timestamp
ms1 = ms;

// Detect edges using Sobel edge detect
sobel(src_int16, dst_int16, NULL);

// Copy timestamp
ms2 = ms;
```

Sobel - algorithm

```
void sobel( const image_t *src, image_t *mag, image_t *dir);
```

See file **EVDK_Operators\spatial_filters.c**

The dir image can be omitted (NULL). If not omitted, it must be of image type **IMGTYPE_FLOAT**.

```
// Scale uint8_pixel_t for convenient visualisation
scaleInt16ToUInt8(dst_int16, dst);

// Convert uint8_pixel_t image to bgr888_pixel_t image for USB
convertToBgr888(dst, usb);
```

Sobel - algorithm

```
void sobel( const image_t *src, image_t *mag, image_t *dir);
```


See file **EVDK_Operators\spatial_filters.c**

The dir image can be omitted (NULL). If not omitted, it must be of image type IMGTYPE_FLOAT.

Optimize most (-O3)
UVC connected

72 ms

EVD1 – Assignment



Study guide
Week 3

5 Spatial filters – sobelFast()

References

- Myler, H. R., & Weeks, A. R. (2009). *The pocket handbook of image processing algorithms in C*. Prentice Hall Press.